



Operating Systems

Deadlocks - Complete

Dr. P. Sateesh Kumar



Resources

- Physical resources
 - like prints, tape drives, memory space, and CPU cycles
- Logical resources
 - like files, semaphores and monitors
- Reusable Resources
 - Processes obtain resources that they later release for reuse by other processes
 - Processors, I/O channels, main and secondary memory, files, databases, and semaphores
- Consumable Resources
 - Created and destroyed by a process
 - Interrupts, signals, messages, and information in I/O buffers



System Model

- Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- Each resource type R_i has W_i instances.
- Each process utilizes a resource as follows:
 - request
 - use
 - release



System Model

Resource-Allocation Graph

A set of vertices V and a set of edges E .

- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system.
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system.
- request edge – directed edge $P_i \rightarrow R_j$
- assignment edge – directed edge $R_j \rightarrow P_i$

System Model

Resource-Allocation Graph

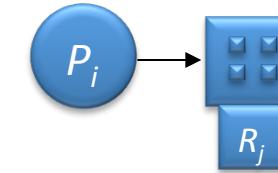
- Process



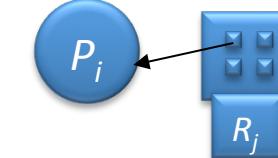
- Resource Type with 4 instances



- P_i requests instance of R_j



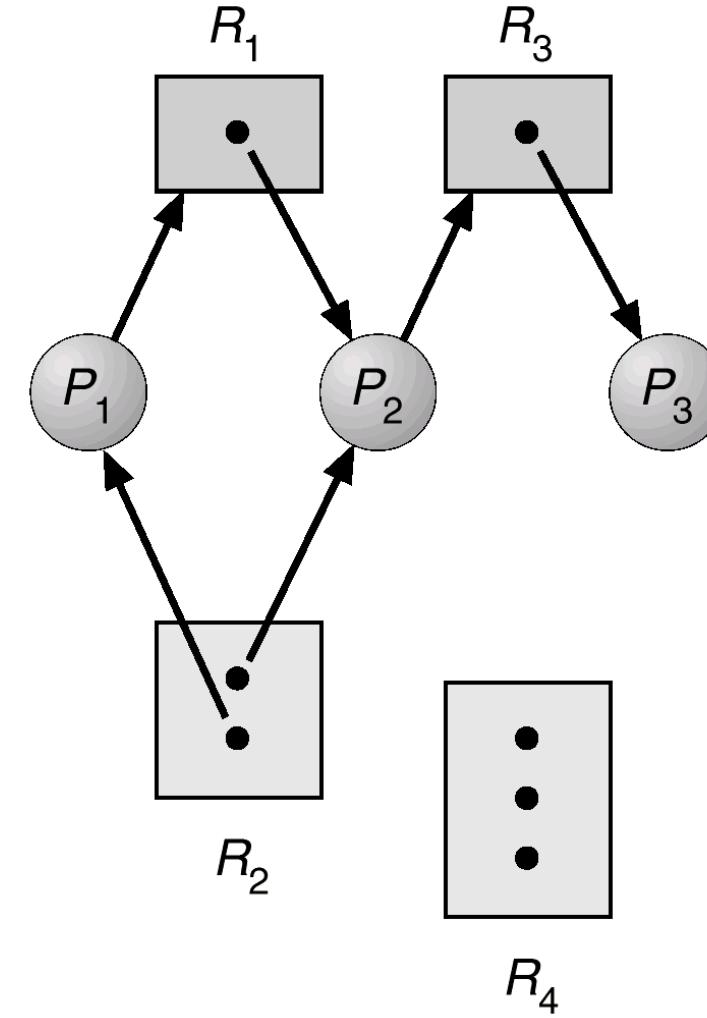
- P_i is holding an instance of R_i





System Model

Example of a Resource Allocation Graph

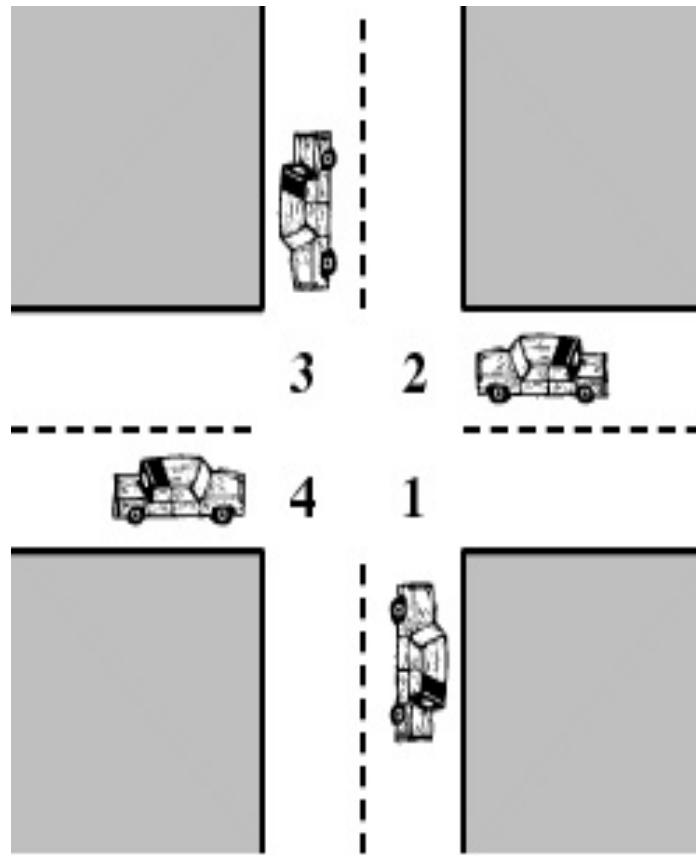




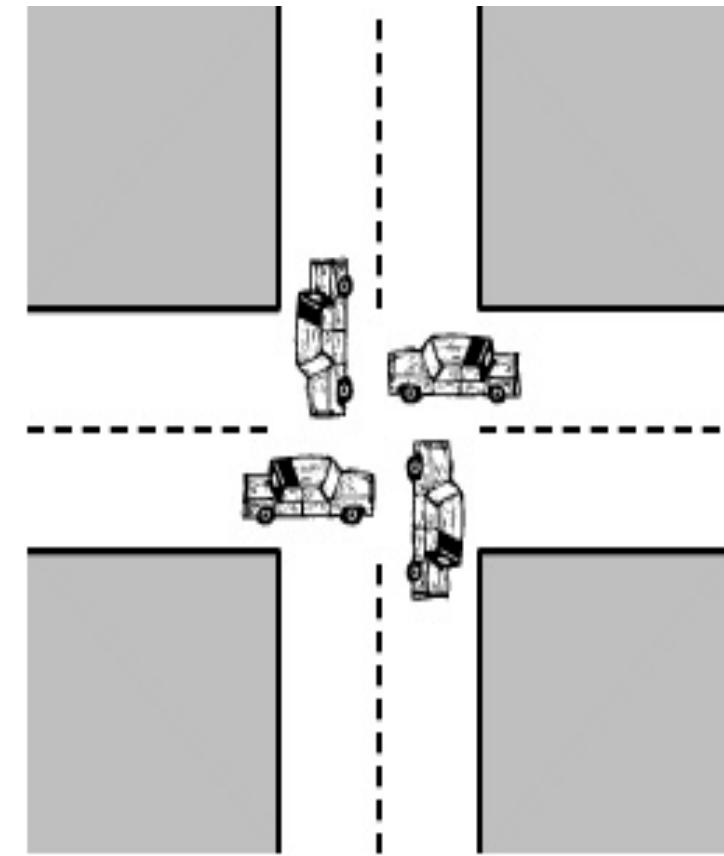
Deadlocks

- In a multiprogramming environment, several processes may compete for a finite number of resources
- A process requests resources, if the resources are not available at that time, the process enters a wait state
- Sometimes a waiting process may never again be able to change state, because the resources they have requested are held by other waiting processes
- This situation is called “**Deadlock**”

Illustration of Deadlock



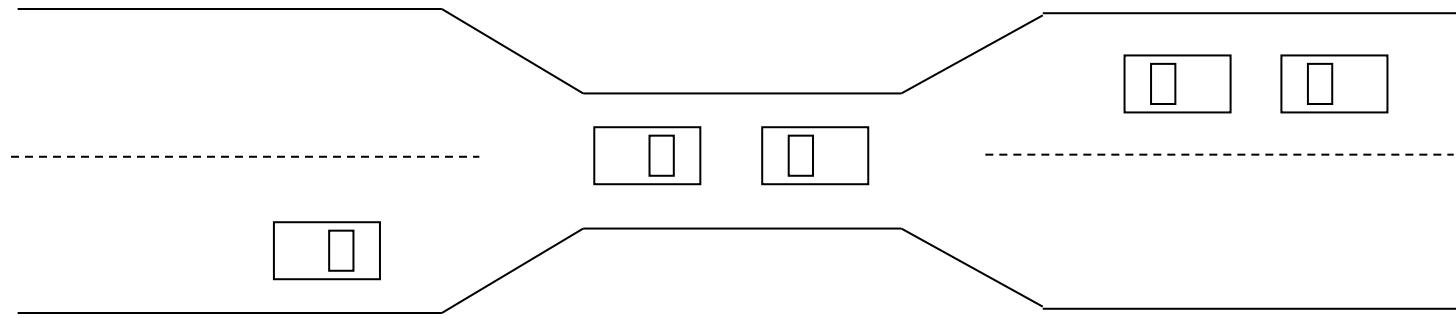
(a) Deadlock possible



(b) Deadlock

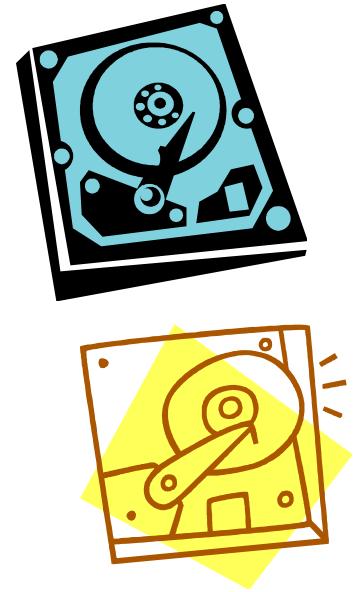


Illustration of Deadlock



Example

- A system has 2 hard disks
- P holds → 1 hard disk (say, A)
- Q holds → 1 hard disk (say, B)
- Both need another hard disk that is held by other process.
- Hence, Deadlock
(if indefinite !!)





Example

Situation - 1

Process P	Process Q
• • •	• • •
Get A	Get B
• • •	• • •
Get B	Get A
• • •	• • •
Release A	Release B
• • •	• • •
Release B	Release A
• • •	• • •

Situation - 2

Process P	Process Q
• • •	• • •
Get A	Get B
• • •	• • •
Release A	Get A
• • •	• • •
Get B	Release A
• • •	• • •
Release B	Release B
• • •	• • •

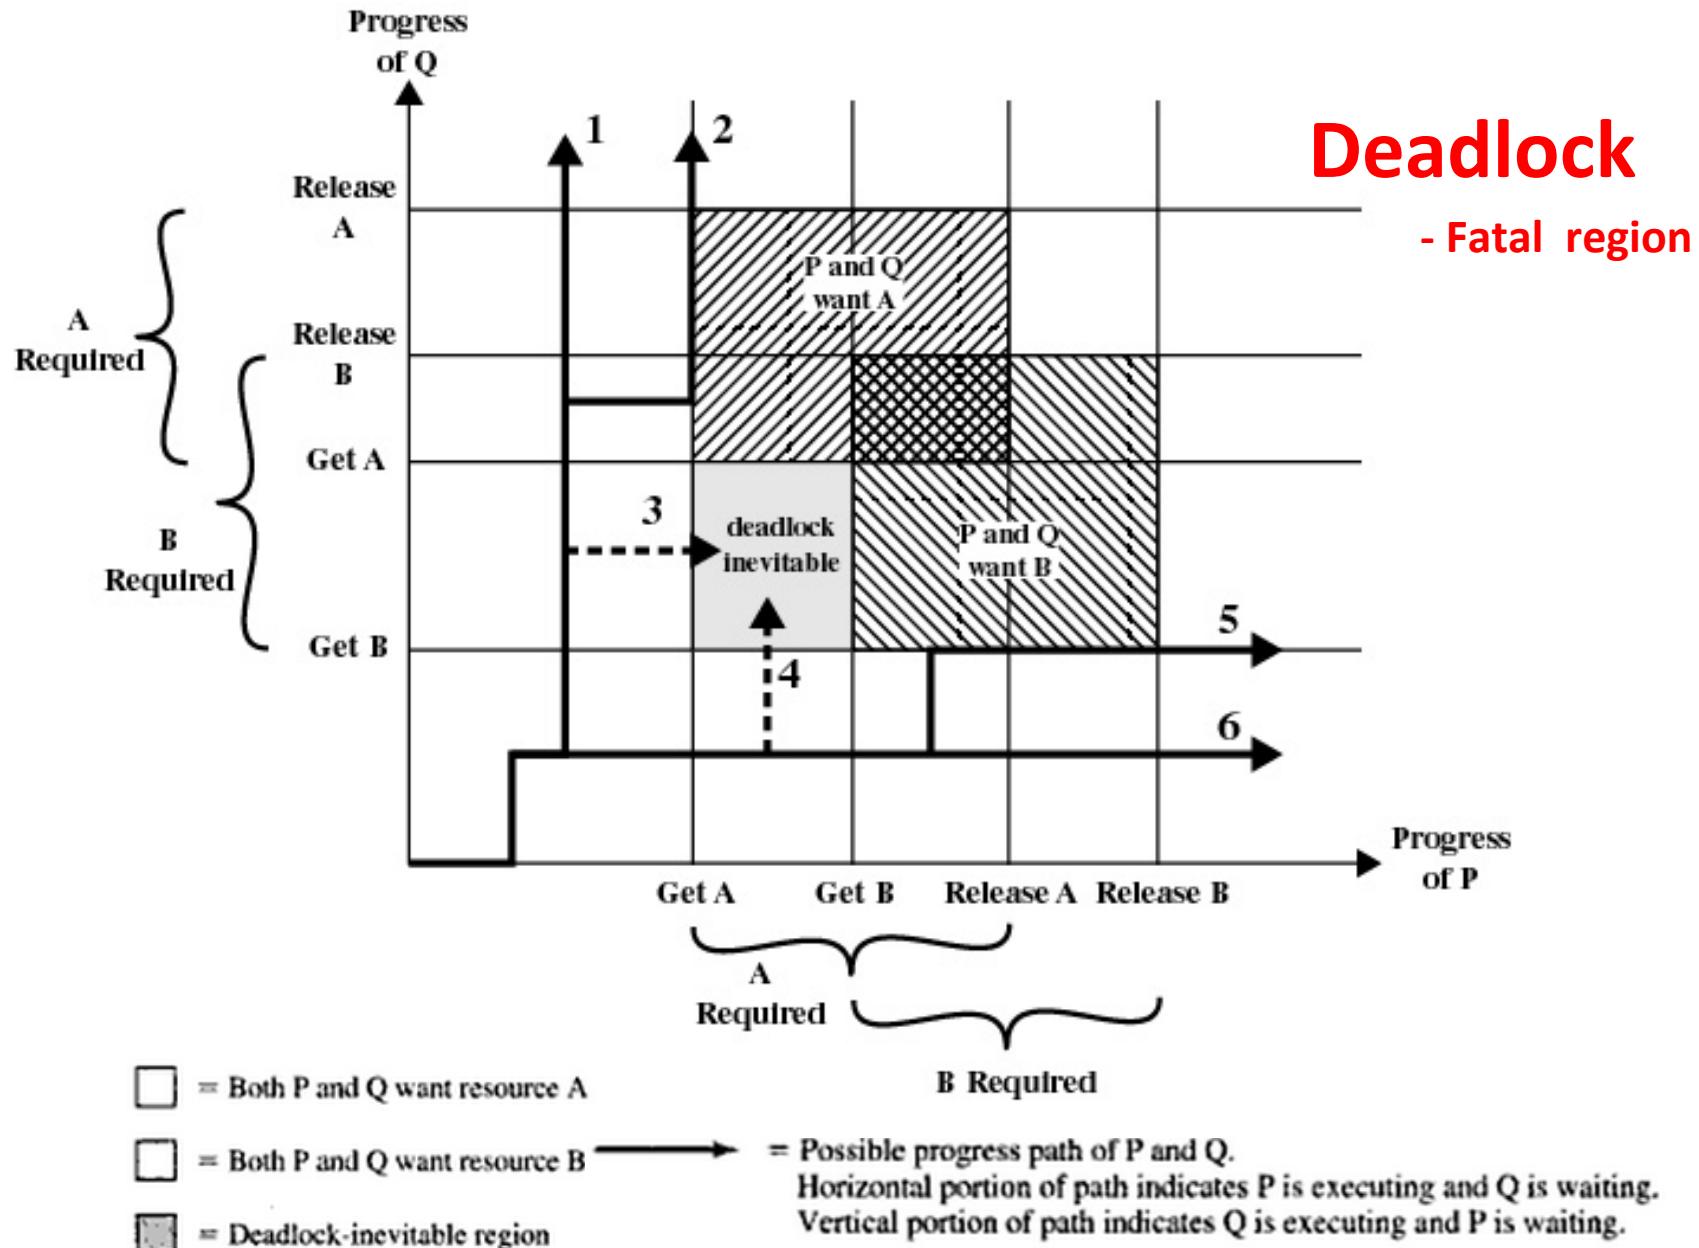


Explanation

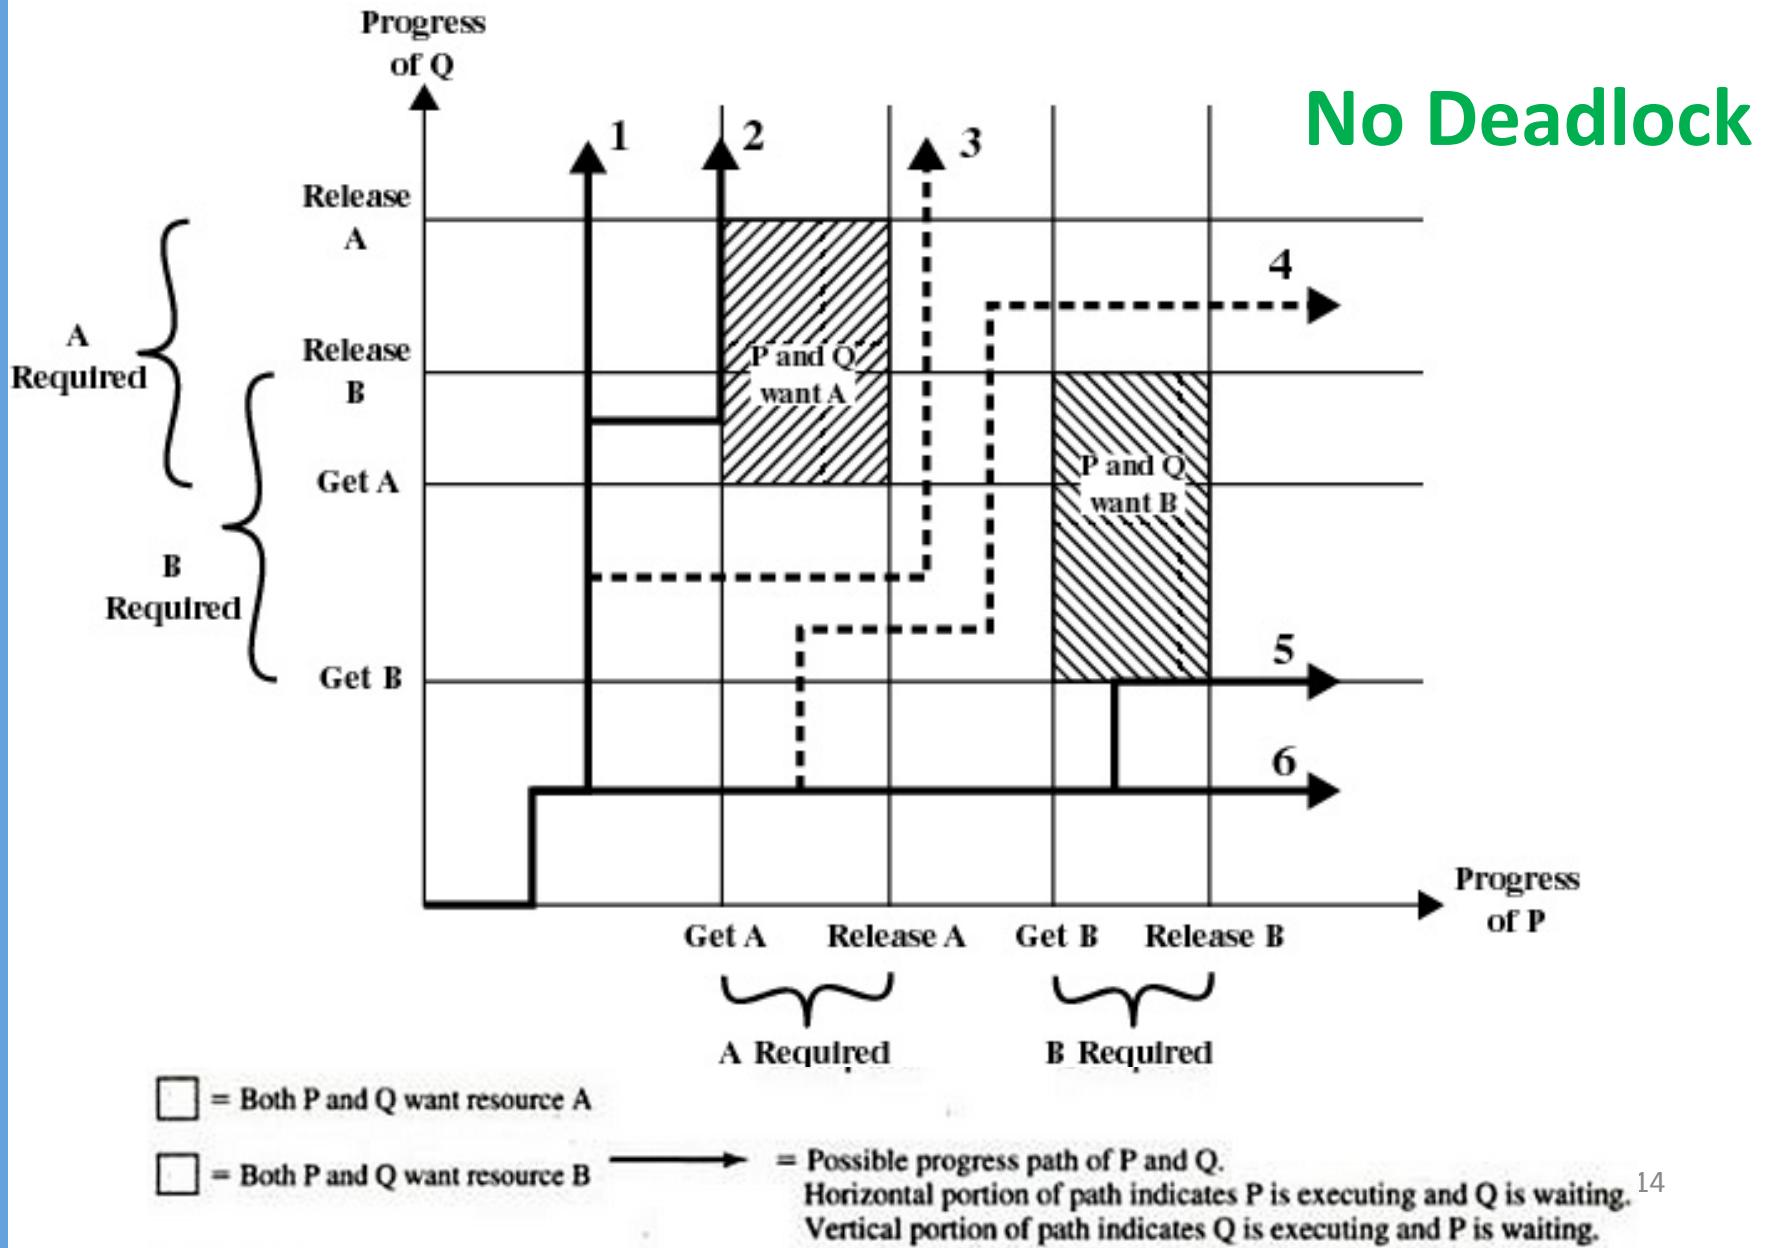
1. Q acquires B and then A and then releases B and A. When P resumes execution, it will be able to acquire both resources.
2. Q acquires B and then A. P executes and blocks on a request for A. Q releases B and A. When P resumes execution, it will be able to acquire both resources.
3. Q acquires B and then P acquires A. Deadlock is inevitable, because as execution proceeds, Q will block on A and P will block on B.
4. P acquires A and then Q acquires B. Deadlock is inevitable, because as execution proceeds, Q will block on A and P will block on B.
5. P acquires A and then B. Q executes and blocks on a request for B. P releases A and B. When Q resumes execution, it will be able to acquire both resources.
6. P acquires A and then B and then releases A and B. When Q resumes execution, it will be able to acquire both resources.



Joint Progress Diagram



Joint Progress Diagram





When will Deadlock Occur?

1. Mutual exclusion
2. Hold and wait
3. No Preemption
4. Circular wait

Necessary Conditions for a Deadlock to Occur



Conditions for Deadlock

- Mutual exclusion
 - only one process may use a resource at a time
- Hold-and-wait
 - A process request all of its required resources at one time



Conditions for Deadlock

- No preemption
 - If a process holding certain resources is denied a further request, that process must release its original resources
 - If a process requests a resource that is currently held by another process, the operating system may preempt the second process and require it to release its resources

Conditions for Deadlock

- Circular wait
 - Prevented by defining a linear ordering of resource types

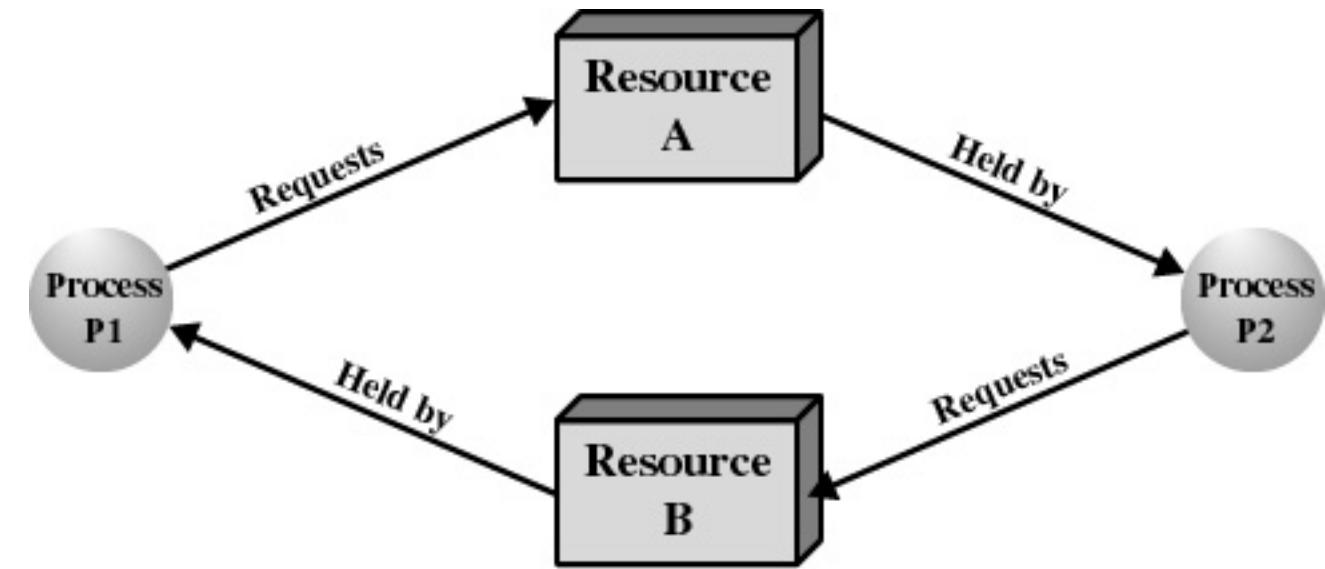
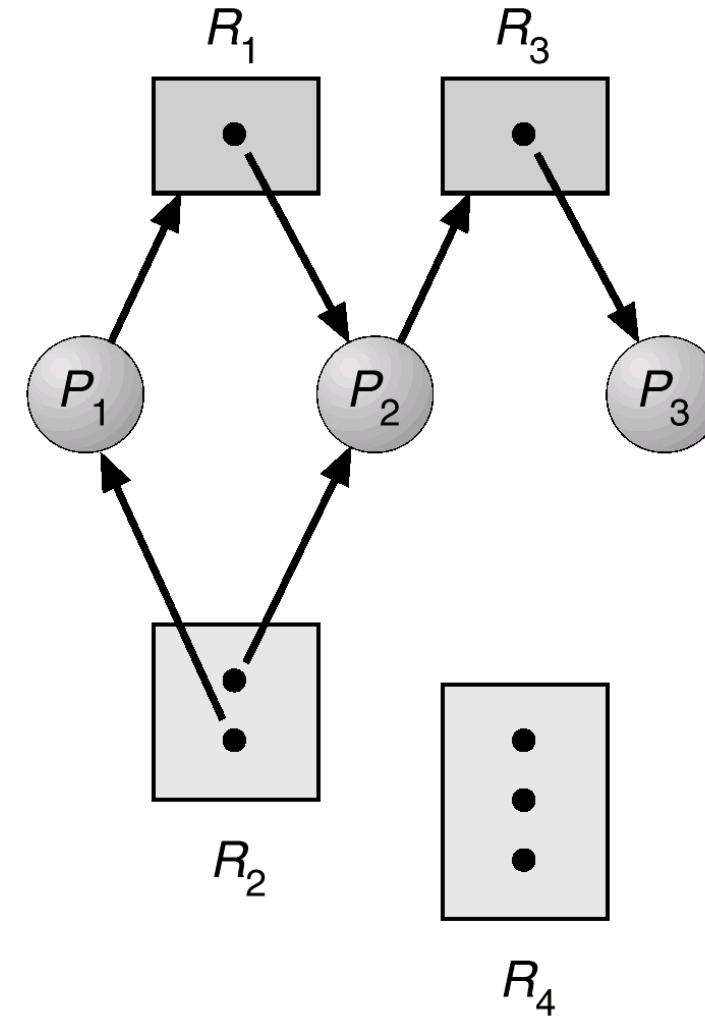
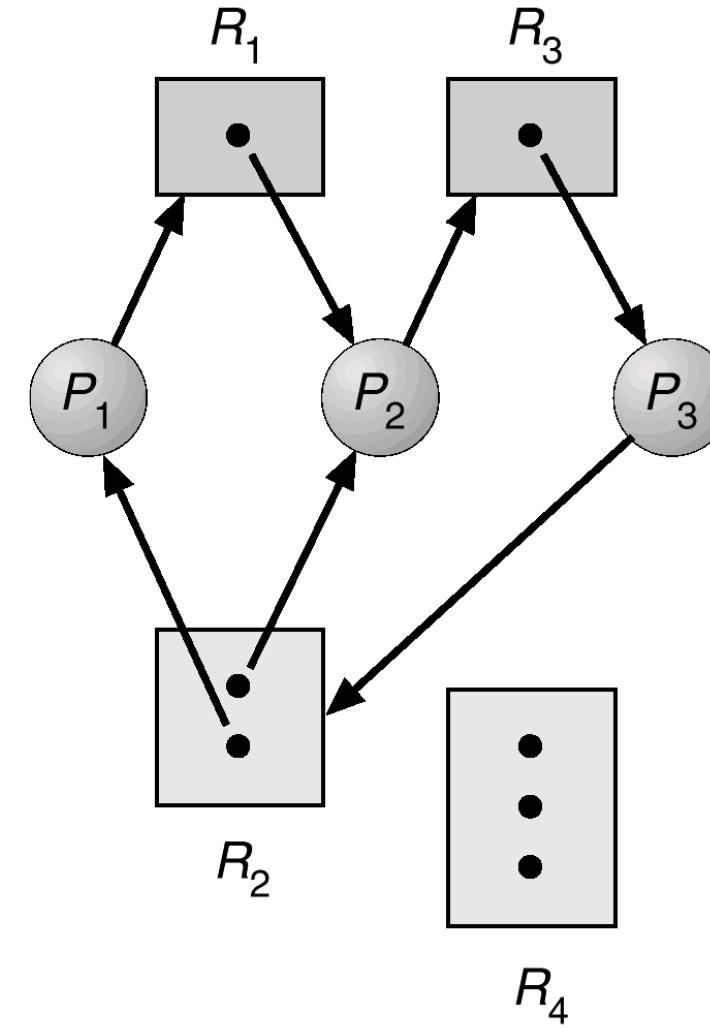


Figure 6.5 Circular Wait

Example - RAG

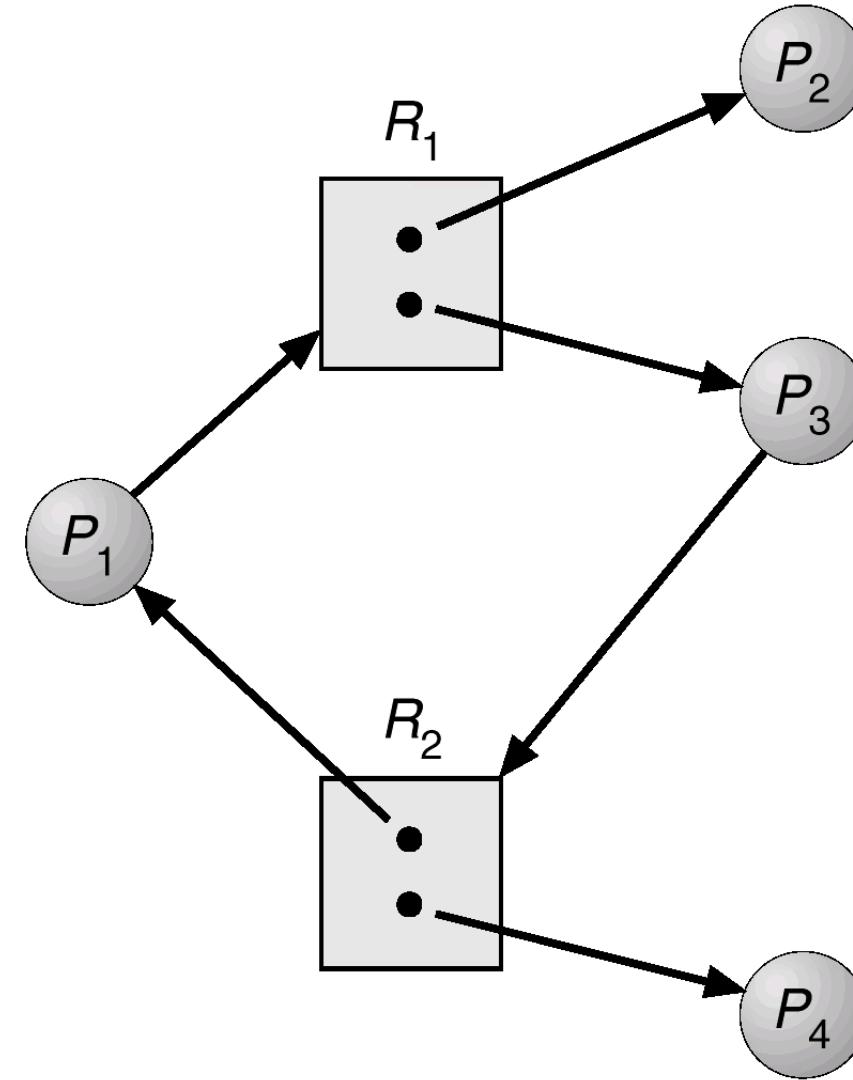


RAG with Deadlock





RAG with Cycle (no Deadlock !)





Basic Facts

- If graph contains no cycles \Rightarrow no deadlock.
- If graph contains a cycle
 - if only one instance per resource type, then deadlock.
 - if several instances per resource type, possibility of deadlock.



Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state. (**Prevention**) (**Avoidance**)
obstacle escape
 - Allow the system to enter a deadlock state and then recover. (**Detection**) (**Recovery**)
 - Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX.

Deadlock Prevention

Restrain the ways request can be made.
(see that at least one of the 4 conditions fail)

- **Mutual Exclusion**
 - not required for sharable resources;
 - Ex: Files for reading (sharable) – allow simultaneous access
 - must hold for non-sharable resources.
 - Printer (non-sharable) – stop simultaneous access



Deadlock Prevention

Restrain the ways request can be made.

- **Hold and Wait**

- must **guarantee** that whenever a process requests a resource, **it does not hold any other resources**.
- Require process to **request** and be allocated **all** its resources **before it begins execution**, or allow process to request resources only when the process has **none**.
- **Low resource utilization; starvation possible.**





Deadlock Prevention



Restrain the ways request can be made.

- **No Preemption**
 - So, preempt the resource
 - If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then **all resources currently being held are released**.
 - Preempted resources are **added to the list of resources for which the process is waiting**.
 - Process will be **restarted only when** it can regain its old resources, as well as the new ones that it is requesting.



Deadlock Prevention

Restrain the ways request can be made.

- **Circular Wait**

- impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration.

$$R = \{R_1, R_2, \dots, R_m\}$$

$R = \{\text{tape drives, disk drives, and printers}\}$

$$F: R \rightarrow N$$

N is a natural number

$$F(\text{tape drive}) = 1$$

$$F(\text{disk drive}) = 5$$

$$F(\text{printer}) = 12$$

$$F(R_j) > F(R_i)$$



Deadlock Avoidance

Requires that the system has some additional *a priori* information available.

- Simplest and most useful model requires that each process **declare the maximum number of resources** of each type that it may **need**.
- The deadlock-avoidance algorithm dynamically **examines the resource-allocation state** to ensure that there can never be a circular-wait condition.
- Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes.



Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a *safe state*.
- System is in safe state if there exists a safe sequence of all processes.
- Sequence $\langle P_1, P_2, \dots, P_n \rangle$ is safe if for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$.



Safe State

- If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished.
- When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate.
- When P_i terminates, P_{i+1} can obtain its needed resources, and so on.



Basic Facts

- If a system is in safe state \Rightarrow no deadlocks.
- If a system is in unsafe state \Rightarrow possibility of deadlock.
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state.

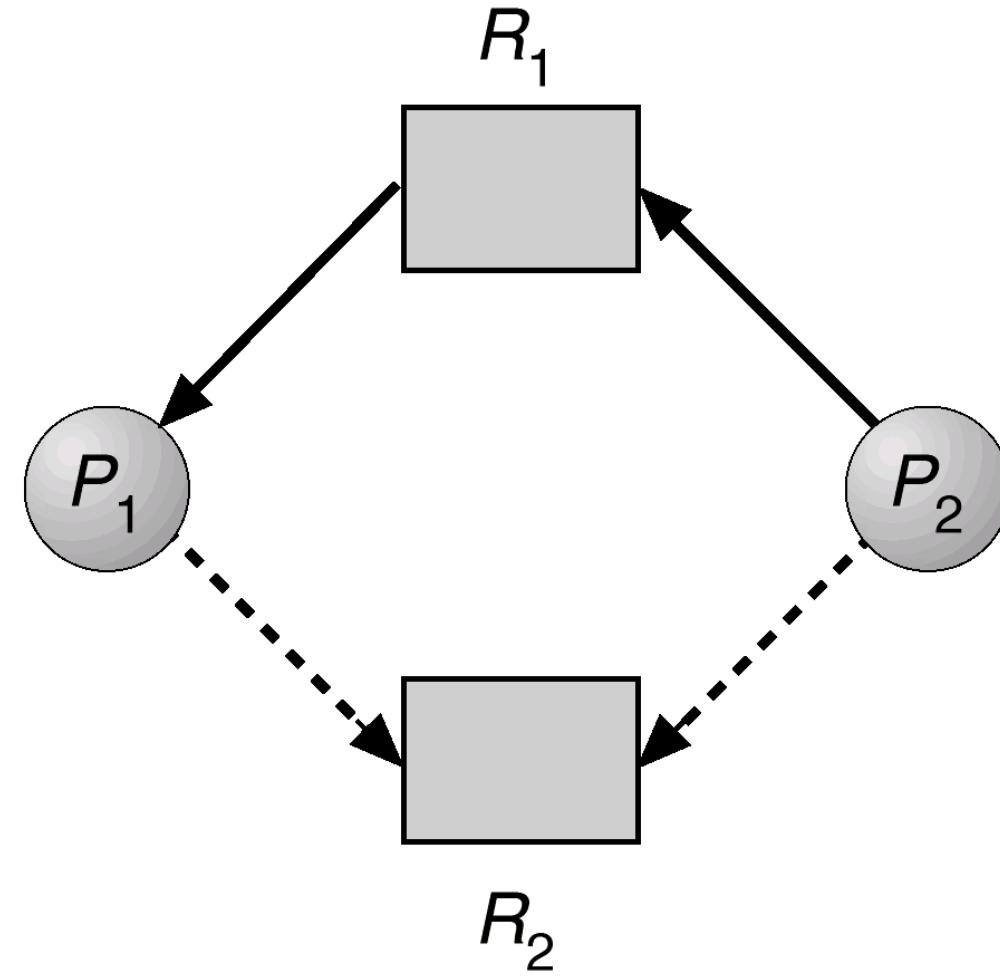


Deadlock Avoidance

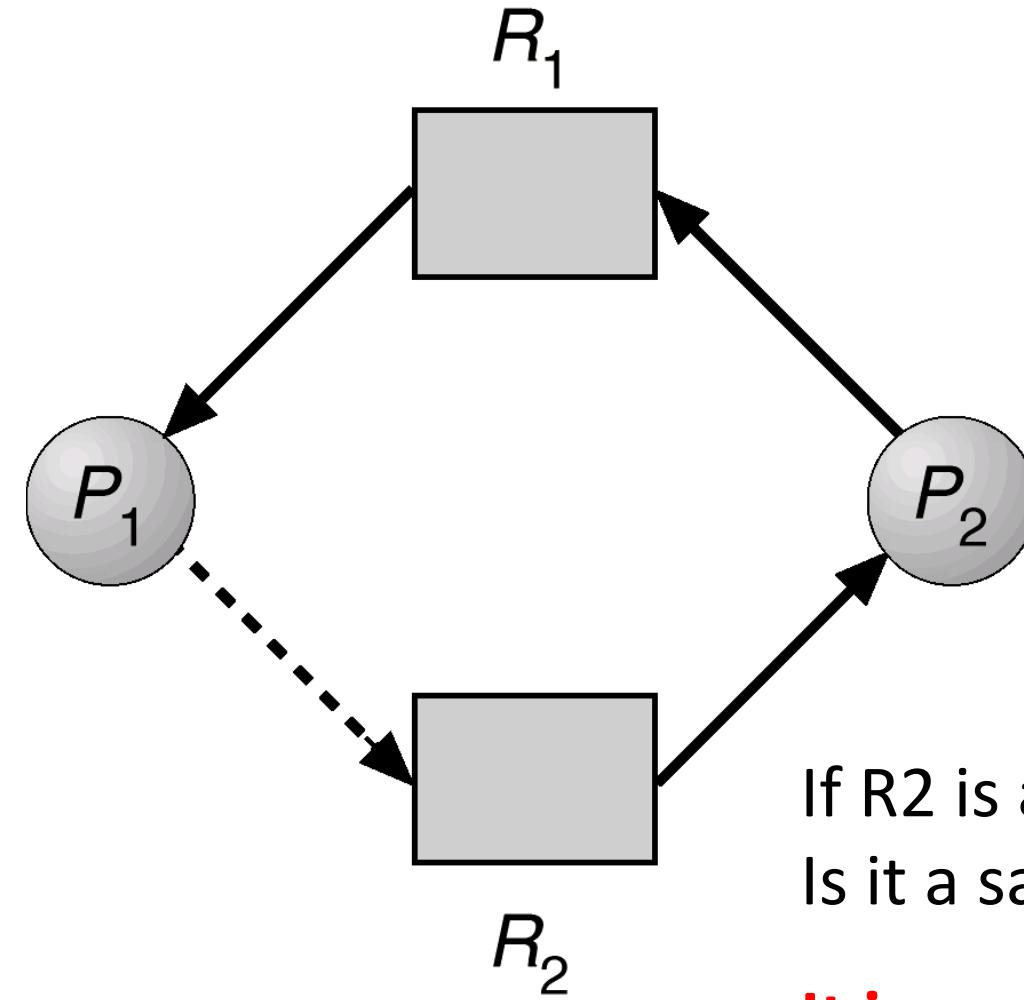
Resource Allocation Graph (RAG)

- ***Claim edge***
 - $P_i \rightarrow R_j$ indicated that process P_j may request resource R_j ; represented by a dashed line.
- Claim edge converts to request edge when a process requests a resource.
- When a resource is released by a process, assignment edge reconverts to a claim edge.
- Resources must be claimed *a priori* in the system.

RAG for Deadlock Avoidance



RAG for Deadlock Avoidance



RAG for Deadlock Avoidance

- Works **only for single instance** of a resource
- It does not work for multiple instances of resources





Illustration of a Situation- Bank

One Resource – Loan (say)

Allocated

Need

Customer	Credit Used	Credit Line
Andy	0	6
Barb	0	5
Marv	0	4
Sue	0	7
Funds Available	10	
Max Commitment		22





Illustration of a Situation- Bank

One Resource – Loan (say)

Allocated

Need

Customer	Credit Used	Credit Line
Andy	1	6
Barb	1	5
Marv	2	4
Sue	4	7
Funds Available	2	
Max Commitment		22

t_1

Illustration of a Situation- Bank

One Resource – Loan (say)

Allocated

Need

Customer	Credit Used	Credit Line
Andy	1	6
Barb	2	5
Marv	2	4
Sue	4	7
Funds Available	1	
Max Commitment		22

t_2



Illustration of a Situation- General

Multiple Resources

Existing

$$\mathbf{E} = (6, 3, 4, 2)$$

Possessed

$$\mathbf{P} = (5, 3, 2, 2)$$

Available

$$\mathbf{A} = (1, 0, 2, 0)$$

$$\mathbf{A} = \mathbf{E} - \mathbf{P}$$

Resources Assigned				
Processes	Tapes	Plotters	Printers	Toasters
A	3	0	1	1
B	0	1	0	0
C	1	1	1	0
D	1	1	0	1
E	0	0	0	0
Total Existing	6	3	4	2
Total Claimed by Processes	5	3	2	2
Remaining Unclaimed	1	0	2	0

Resources Still Needed				
Processes	Tapes	Plotters	Printers	Toasters
A	1	1	0	0
B	0	1	1	2
C	3	1	0	0
D	0	0	1	0
E	2	1	1	0

Illustration of a Situation- General

1. Look for a row whose unmet needs don't exceed what's available. i.e., a row where $P \leq A$ if no such row exists, the system is **deadlocked** because no process can acquire the resources it needs to run to completion. If there's more than one such row, just **pick one**
2. Assume that the process chosen in 1 **acquires** all the resources it needs and runs to completion, thereby **releasing** its resources. **Mark that process** as **virtually terminated** and **add** its resources to A.
3. Repeat 1 and 2 until all processes are either virtually terminated (**safe state**), or a deadlock is detected (**unsafe state**)

Resources Assigned				
Processes	Tapes	Plotters	Printers	Toasters
A	3	0	1	1
B	0	1	0	0
C	1	1	1	0
D	1	1	0	1
E	0	0	0	0
Total Existing	6	3	4	2
Total Claimed by Processes	5	3	2	2
Remaining Unclaimed	1	0	2	0

Resources Still Needed				
Processes	Tapes	Plotters	Printers	Toasters
A	1	1	0	0
B	0	1	1	2
C	3	1	0	0
D	0	0	1	0
E	2	1	1	0

Illustration of a Situation- General

- As per the Algorithm, the process D's requirements are smaller than A.
 $E = \{6, 3, 4, 2\}$
 $P = \{5, 3, 2, 2\} - \{1, 1, 0, 1\} = \{4, 2, 2, 1\}$
 $A = \{1, 0, 2, 0\} + \{1, 1, 0, 1\} = \{2, 1, 2, 1\}$
- Now, A's requirements are less than A. So, do the same thing with A:
 $P = \{4, 2, 2, 1\} - \{3, 0, 1, 1\} = \{1, 2, 1, 0\}$
 $A = \{2, 1, 2, 1\} + \{3, 0, 1, 1\} = \{5, 1, 3, 2\}$
- At this point, we see that there are no remaining processes that can't be satisfied from available resources. So the illustrated state is **safe**.

Resources Assigned				
Processes	Tapes	Plotters	Printers	Toasters
A	3	0	1	1
B	0	1	0	0
C	1	1	1	0
D	1	1	0	1
E	0	0	0	0
Total Existing	6	3	4	2
Total Claimed by Processes	5	3	2	2
Remaining Unclaimed	1	0	2	0

Resources Still Needed				
Processes	Tapes	Plotters	Printers	Toasters
A	1	1	0	0
B	0	1	1	2
C	3	1	0	0
D	0	0	1	0
E	2	1	1	0



Banker's Algorithm for Deadlock Avoidance

- Multiple instances.
- Each process must *a priori* claim maximum use.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.



Banker's Algorithm for Deadlock Avoidance

- Resource – Request Algorithm
 - Verifies if the requested resource can be safely granted.
- Safety Algorithm
 - To check whether or not a system is in a safe state.



Banker's Algorithm for Deadlock Avoidance

Data Structures:

Let n = number of processes, and m = number of resources types.

- **Available**: Vector of length m .

If $\text{available}[j] = k$, there are k instances of resource type R_j available.

- **Allocation**: $n \times m$ matrix.

If $\text{Allocation}[i,j] = k$ then P_i is currently allocated k instances of R_j .

- **Max**: $n \times m$ matrix.

If $\text{Max}[i,j] = k$, then process P_i may request at most k instances of resource type R_j .

- **Need**: $n \times m$ matrix.

If $\text{Need}[i,j] = k$, then P_i may need k more instances of R_j to complete its task.

Available				Allocation				Max		
A	B	C		A	B	C	A	B	C	
3	3	2		P ₀	0	1	0	7	5	3
				P ₁	2	0	0	3	2	2
				P ₂	3	0	2	9	0	2
				P ₃	2	1	1	2	2	2
				P ₄	0	0	2	4	3	3

$$\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j].$$



Banker's Algorithm for Deadlock Avoidance

Resource Request Algorithm

Request = request vector for process P_i ,

If $Request_i[j] = k$ then process P_i wants k instances of resource type R_j .

1. If $Request_i \leq Need_i$, go to step 2 (error?). Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If $Request_i \leq Available$, go to step 3 (min?). Otherwise P_i must wait, since resources are not available.
3. Pretend to allocate requested resources to P_i by modifying the state as follows:

(safe?) $Available = Available - Request_i;$

$Allocation_i = Allocation_i + Request_i;$

$Need_i = Need_i - Request_i;$

• If safe \Rightarrow the resources are allocated to P_i . (Run the Safety Algorithm)

• If unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored

Available		Allocation	Max	Need
A B C		A B C	A B C	A B C
3 3 2	P0	0 1 0	7 5 3	7 4 3
	P1	2 0 0	3 2 2	1 2 2
	P2	3 0 2	9 0 2	6 0 0
	P3	2 1 1	2 2 2	0 1 1
	P4	0 0 2	4 3 3	4 3 1



Banker's Algorithm for Deadlock Avoidance

Safety Algorithm

1. Let $Work$ and $Finish$ be vectors of length m and n , respectively.
Initialize:

$m = \text{No. of Resources}$ $n = \text{No. of Processes}$

$Work = Available$

$Finish [i] = \text{false}$ for $i = 1, 2, 3, \dots, n$.

2. Find an i (process) such that **both**:

(a) $Finish [i] = \text{false}$

(b) $Need_i \leq Work$

If no such i exists, go to step 4.

3. $Work = Work + Allocation_i$

$Finish[i] = \text{true}$

go to step 2.

4. If $Finish [i] == \text{true}$ for all i , then the system is in a safe state.

Available	Allocation	Max	Need
A B C	A B C	A B C	A B C
3 3 2	P0 0 1 0	7 5 3	7 4 3
	P1 2 0 0	3 2 2	1 2 2
	P2 3 0 2	9 0 2	6 0 0
	P3 2 1 1	2 2 2	0 1 1
	P4 0 0 2	4 3 3	4 3 1



Example of Banker's Algorithm

- 5 processes P_0 through P_4 ;
- 3 resource types A (10 instances), B (5 instances), and C (7 instances).

0 2 0

Available		Allocation	Max	Need								
A	B	C	A	B	C	A	B	C				
3	3	2	P_0	0	3	0	7	5	3	7	2	3
3 2 0			P_1	2	0	0	3	2	2	1	2	2
1 0 5			P_2	3	0	2	9	0	2	6	0	0
1 0 5			P_3	2	1	1	2	2	2	0	1	1
1 0 5			P_4	0	0	2	4	3	3	4	3	1

Snapshot at time T_0

Safe? $\langle P_1, P_3, P_4, P_2, P_0 \rangle$

Example of Banker's Algorithm

Suppose P_1 now requests(1,0,2)

Available = (3,3,2)

Check that Request \leq Available ? $(1,0,2) \leq (3,3,2) \Rightarrow \text{true.}$

Available		Allocation	Max	Need					
A	B	C	A	B	C	A	B	C	
2	3	0	P ₀	0	3	0	7	5	3
3	3	2	P ₁	3	0	2	7	2	3
			P ₂	3	0	2	3	2	2
			P ₃	2	1	1	9	0	2
			P ₄	0	0	2	2	2	6
							0	1	1
							4	3	1

- Executing safety algorithm

shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement.

- Can request for (3,3,0) by P_4 be granted?
- Can request for (0,2,0) by P_0 be granted?